

Empirical Bayes shrinkage reduces variance, but introduces bias

Red line is $x = y$; dashed white line is a linear fit

